

Ahmed Ramzan

AI Engineer

Sahiwal Division, Punjab, Pakistan — 0322 8968216 — ahmadramzan2219@gmail.com — linkedin.com/in/ahmed-r-0993ba253
— github.com/ahmadramzan9295

Summary

AI Engineer with 1 year of professional experience building Generative AI, LLM-powered applications, RAG systems, AI agents, computer vision solutions, and production-ready APIs. Experienced in integrating OpenAI, Anthropic, OpenRouter, and Ollama models, developing Python/FastAPI services, designing retrieval and prompt workflows, and deploying containerized applications using Docker and AWS. Strong foundation in machine learning, deep learning, computer vision

Skills

- **Languages:** Python, JavaScript, SQL, R
- **Generative AI / LLM:** Prompt Engineering, AI Agents, RAG, Embeddings, Semantic Search, Structured Outputs, LLM API Integration (OpenAI, Anthropic, OpenRouter, Ollama)
- **ML / DL:** Scikit-learn, PyTorch, TensorFlow, Keras, Computer Vision, NLP, Model Evaluation
- **Backend:** FastAPI, Flask, REST APIs, SQLAlchemy, OAuth 2.0
- **Data:** Pandas, NumPy, EDA, Feature Engineering
- **Deployment:** Docker, AWS, Git, GitHub

Experience

AI Engineer

Dec 2025 – August 2026

AibitSoft Solution

- Develop and deliver production-oriented AI solutions for international clients across Generative AI, LLM applications, computer vision, NLP, and predictive analytics.
- Build RAG applications using document ingestion, intelligent chunking, embeddings, vector search, retrieval pipelines, and LLM-based response generation to ground answers in user-provided knowledge.
- Develop LLM-powered AI agents and automation workflows using Python, prompt engineering, tool integration, structured outputs, and multi-step task execution.
- Integrate and evaluate LLM providers including OpenAI, Anthropic, OpenRouter, and Ollama, selecting models and designing prompts based on application requirements.
- Build backend services and REST APIs with FastAPI and Flask for AI inference, model integration, authentication, data processing, and application workflows.
- Develop machine learning and deep learning pipelines using Scikit-learn, PyTorch, TensorFlow, Pandas, and NumPy for classification, prediction, NLP, and computer vision tasks.
- Containerize AI applications with Docker and prepare deployment workflows for cloud environments including AWS, covering application packaging, configuration, inference testing, and service integration.

Generative AI Intern

Sep 2025 – Oct 2025

Arch Technologies – Remote, Pakistan

- Built and evaluated Generative AI applications and task-automation workflows using LLM APIs, prompt engineering, and Python-based application logic.
- Integrated LLM APIs with structured JSON inputs/outputs and improved application reliability through prompt optimization, testing, and debugging.

Projects

ROTSAM – Referring Remote Sensing Image Segmentation (Research)

Torch, SAM3, Vision Transformers, Optimal Transport

- Designed a referring remote sensing image segmentation framework built on frozen SAM3 vision and text encoders for pixel-level segmentation of objects described by natural-language expressions.
- Developed a Dynamic Low-Rank Vision Transformer and Multi-Scale Optimal Transport alignment module to address scale variation, orientation, and background clutter.

ApexScan POS – Smart Hardware-Integrated POS & Package Dimensioning System

Python, PyQt6, OpenCV, Docker, SQLite, PySerial, HIDAPI, Google Cloud Vision OCR

- Built a multi-threaded PyQt6 desktop POS app integrating scales and barcode scanners via low-level pyserial/hidapi/pyusb drivers, with an OpenCV + Google Cloud Vision OCR pipeline for real-time parcel dimensioning and carrier ID; containerized with Docker (USB/video passthrough) and synced to a cloud backend via REST APIs over local SQLite storage.

Education

Bachelor of Science in Data Science

2022 – 2026

GIFT University

Relevant coursework: Machine Learning, Deep Learning, Natural Language Processing, Computer Vision